

# solutions\_R

May 19, 2022

## 1 ASRR Messy Data Challenge

### 1.1 Example analysis (R version)

```
[ ]: # Load library
packages <- c("tidyverse", "data.table", "haven", "skimr", "janitor")
for (p in packages) {
  if (!require(p, character.only = TRUE)) install.packages(p)
  suppressPackageStartupMessages(library(p, character.only = TRUE))
}
```

### 1.2 Data exploration

#### 1.2.1 Read in data

```
[2]: df <- read_dta("../data/icu_data.dta")
head(df)
```

A tibble: 6 × 13

	age <dbl>	gender <dbl>	iculos <dbl>	hr <dbl>	temp <dbl>	sbp <dbl>	dbp <dbl>	resp <dbl>	o2sat <dbl>	map <dbl>	sep <dbl>
	65.36	0	28	71.0	38.11	168.00	56.0	17	94.0	91.0	0
	55.00	0	7	76.0	37.50	133.00	74.0	24	96.0	99.0	0
	38.00	0	15	77.0	NA	146.00	83.0	17	NA	108.0	0
	72.17	1	34	53.5	NA	157.00	60.5	15	95.5	92.5	0
	61.26	1	2	72.5	NA	93.25	68.5	18	100.0	79.0	0
	21.00	1	19	74.0	NA	118.00	104.0	20	91.0	112.0	0

#### 1.2.2 What's in the dataset

```
[3]: summary(df)
```

age	gender	iculos	hr
Min. : 14.00	Min. : 0.0000	Min. : 1.00	Min. : 20.00
1st Qu.: 51.00	1st Qu.: 0.0000	1st Qu.: 10.00	1st Qu.: 72.00
Median : 63.44	Median : 1.0000	Median : 21.00	Median : 83.00
Mean : 61.74	Mean : 0.5485	Mean : 26.18	Mean : 84.23
3rd Qu.: 74.00	3rd Qu.: 1.0000	3rd Qu.: 34.00	3rd Qu.: 95.50
Max. : 100.00	Max. : 1.0000	Max. : 336.00	Max. : 223.00
			NA's : 131167

temp	sbp	dbp	resp
Min. :20.9	Min. : 20.0	Min. : 20	Min. : 1.00
1st Qu.:36.5	1st Qu.:107.0	1st Qu.: 55	1st Qu.: 16.00
Median :36.9	Median :122.0	Median : 63	Median : 18.00
Mean :36.9	Mean :124.7	Mean : 65	Mean : 18.73
3rd Qu.:37.4	3rd Qu.:140.0	3rd Qu.: 73	3rd Qu.: 21.00
Max. :50.0	Max. :298.0	Max. :300	Max. :100.00
NA's :822321	NA's :176747	NA's :373178	NA's :205042

o2sat	map	sepsislabel	hospid
Min. : 20.00	Min. : 20.00	Min. :0.000000	Length:1201974
1st Qu.: 96.00	1st Qu.: 72.00	1st Qu.:0.000000	Class :character
Median : 98.00	Median : 82.00	Median :0.000000	Mode :character
Mean : 97.15	Mean : 83.55	Mean :0.001379	
3rd Qu.: 99.00	3rd Qu.: 93.00	3rd Qu.:0.000000	
Max. :100.00	Max. :300.00	Max. :1.000000	
NA's :170544	NA's :163352		

patid

Min. : 1

1st Qu.:12317

Median :23294

Mean :22012

3rd Qu.:31822

Max. :40336

### 1.2.3 Distributions of each of the variables

```
[4]: options(width = 110)
skim(df)
```

```
Data Summary
Name      Values
Number of rows      1201974
Number of columns    13

-----
Column type frequency:
  character          1
  numeric            12

-----
Group variables      None

Variable type: character

skim_variable n_missing complete_rate  min  max empty n_unique whitespace
1 hospid      0              1      1    1    0        2          0
```

Variable type: numeric

	skim_variable	n_missing	complete_rate		mean	sd	p0	p25
p50	p75	p100	hist					
1	age		0	1	61.7	16.5	14	51
63.4	74	100						
2	gender		0	1	0.549	0.498	0	0
1	1	1						
3	iculos		0	1	26.2	27.9	1	10
21	34	336						
4	hr		131167	0.891	84.2	17.6	20	72
83	95.5	223						
5	temp		822321	0.316	36.9	0.759	20.9	36.5
36.9	37.4	50						
6	sbp		176747	0.853	125.	23.6	20	107
122	140	298						
7	dbp		373178	0.690	65.0	14.2	20	55
63	73	300						
8	resp		205042	0.829	18.7	5.02	1	16
18	21	100						
9	o2sat		170544	0.858	97.1	2.98	20	96
98	99	100						
10	map		163352	0.864	83.5	16.6	20	72
82	93	300						
11	sepsislabel		0	1	0.00138	0.0371	0	0
0	0	1						
12	patid		0	1	22012.	11502.	1	12317
23294	31822	40336						

#### 1.2.4 Complete case indicator

```
[5]: cols <- c("o2sat", "hr", "temp", "sbp", "map", "resp")
df <- df %>%
  mutate(nvar_miss = rowSums(across(all_of(cols), is.na)),
         cc_fl = ifelse(nvar_miss == 0, TRUE, FALSE))
```

```
[6]: tabyl(df, cc_fl)
```

	cc_fl	n	percent
	<lgl>	<int>	<dbl>
A tabyl: 2 × 3	FALSE	865835	0.7203442
	TRUE	336139	0.2796558

Only 28% of records have no missing vital signs

## 1.3 Outcome exploration

### 1.3.1 How many people were diagnosed with sepsis?

```
[7]: df %>%  
      group_by(patid) %>%  
      summarise(any_sepsis = max(sepsislabel)) %>%  
      tabyl(any_sepsis)
```

	any_sepsis	n	percent
	<dbl>	<int>	<dbl>
A tabyl: 2 × 3	0	29268	0.94641876
	1	1657	0.05358124

### 1.3.2 When do people get sepsis in ICU?

```
[8]: df_sepsis <- df %>%  
      filter(sepsislabel == 1) %>%  
      group_by(patid) %>%  
      summarise(time_to_sepsis = min(iculos))  
      skim(df_sepsis, time_to_sepsis)
```

#### Data Summary

Name	Values
df_sepsis	
Number of rows	1657
Number of columns	2

#### Column type frequency:

numeric	1
---------	---


Group variables	None
-----------------	------

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100
time_to_sepsis	0	1	65.7	60.9	7	20	45	91	331

- min: 7 hours
- max: 331 hours (13.8 days)
- median: 45 hours

```
[9]: ggplot(df_sepsis, aes(x = time_to_sepsis)) +  
      theme_bw() +  
      # these settings are used to match stata output  
      geom_histogram(fill = "#2a9d8f", colour = "black", bins = 32,  
                     binwidth = 10.125, boundary = 7)
```



### 1.3.3 Create indicator for patient who get sepsis:

```
[10]: df <- df %>%  
      group_by(patid) %>%  
      mutate(any_sepsis = max(sepsislabel, na.rm = TRUE))
```

## 1.4 Imputing explanatory measures

### 1.4.1 Mean Imputation

```
[11]: impute_mean <- function(x) replace_na(x, mean(x, na.rm=T))
# Use only ICULOS <= 5
df_imp <- df %>%
  filter(iculos <= 5) %>%
  group_by(patid) %>%
  arrange(patid, iculos) %>%
  mutate(across(all_of(cols), impute_mean, .names = "{.col}_imp1"))
```

### 1.4.2 First observation carried backwards

```
[12]: df_imp <- df_imp %>%
  mutate(across(all_of(cols), ~.x, .names = "{.col}_imp2")) %>%
  fill(ends_with("_imp2"), .direction = "up")

head(df_imp)
```

	age	gender	iculos	hr	temp	sbp	dbp	resp	o2sat	map
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
A grouped_df: 6 × 28	83.14	0	1	NA	NA	NA	NA	NA	NA	NA
	83.14	0	2	97	NA	98	NA	19.0	95.0	75.33
	83.14	0	3	89	NA	122	NA	22.0	99.0	86.00
	83.14	0	4	90	NA	NA	NA	30.0	95.0	NA
	83.14	0	5	103	NA	122	NA	24.5	88.5	91.33
	75.91	0	1	NA	NA	NA	NA	NA	NA	NA

### 1.4.3 Inspect missingness again among imputed variables

```
[13]: df_imp <- df_imp %>%
  ungroup() %>%
  mutate(nvar_miss_imp1 = rowSums(across(ends_with("_imp1"), is.na)),
         nvar_miss_imp2 = rowSums(across(ends_with("_imp2"), is.na)),
         cc_fl_imp1 = ifelse(nvar_miss_imp1 == 0, TRUE, FALSE),
         cc_fl_imp2 = ifelse(nvar_miss_imp2 == 0, TRUE, FALSE))
```

```
[14]: df_imp %>%
  filter(iculos == 1) %>%
  tabyl(cc_fl_imp1)
```

	cc_fl_imp1	n	percent
	<lgl>	<int>	<dbl>
A tabyl: 2 × 3	FALSE	6897	0.2230234
	TRUE	24028	0.7769766

```
[15]: df_imp %>%
      filter(iculos == 1) %>%
      tabyl(cc_fl_imp2)
```

	cc_fl_imp2	n	percent
	<lgl>	<int>	<dbl>
A tabyl: 2 × 3	FALSE	6897	0.2230234
	TRUE	24028	0.7769766

78% of rows non-missing for each imputation method

## 1.5 Modelling

### 1.5.1 Dummy indicators for hospital:

In R, dummy indicators for a binary / categorical (a.k.a factor) variable will be created automatically when creating a formula object to be used in generalised-linear model with `glm()`.

However, a no-intercept model is needed when fitting the regression to avoid multicollinearity issue due to singular matrix. For more discussion, see: <https://stats.stackexchange.com/a/94021>

### 1.5.2 Mean imputation

```
[16]: model_imp1 <- glm(
  # the -1 term will suppress intercept, which is needed to avoid singularity
  ↪matrix issue
  any_sepsis ~ -1 + age + gender + o2sat_imp1 + hr_imp1 + temp_imp1
               + sbp_imp1 + map_imp1 + resp_imp1 + factor(hospid),
  data = filter(df_imp, iculos == 1),
  family = binomial(link = "logit")
)
```

```
[17]: summary(model_imp1)
```

Call:

```
glm(formula = any_sepsis ~ -1 + age + gender + o2sat_imp1 + hr_imp1 +
    temp_imp1 + sbp_imp1 + map_imp1 + resp_imp1 + factor(hospid),
    family = binomial(link = "logit"), data = filter(df_imp,
    iculos == 1))
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.0167	-0.3480	-0.2962	-0.2514	2.9621

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
age	0.001538	0.001873	0.821	0.41162
gender	0.182319	0.060440	3.017	0.00256 **
o2sat_imp1	0.014192	0.011535	1.230	0.21855

```

hr_imp1      0.011620    0.001835    6.334 2.39e-10 ***
temp_imp1    0.007815    0.040582    0.193 0.84730
sbp_imp1     0.002920    0.002554    1.143 0.25302
map_imp1     -0.016736    0.003827   -4.373 1.23e-05 ***
resp_imp1    0.050361    0.006298    7.996 1.29e-15 ***
factor(hospid)A -5.496563  1.878954   -2.925 0.00344 **
factor(hospid)B -5.923873  1.876972   -3.156 0.00160 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 33309.9 on 24028 degrees of freedom
Residual deviance: 9321.4 on 24018 degrees of freedom
(6897 observations deleted due to missingness)
AIC: 9341.4

```

Number of Fisher Scoring iterations: 6

```

[18]: # Calculate odds ratio & 95% confidence interval
exp(coefficients(model_imp1)) %>%
  enframe(name = "variable", value = "odds ratio") %>%
  add_column(as_tibble(exp(confint(model_imp1))))

```

Waiting for profiling to be done...

	variable <chr>	odds ratio <dbl>	2.5 % <dbl>	97.5 % <dbl>
	age	1.00153917	9.978819e-01	1.0052367
	gender	1.19999753	1.066285e+00	1.3514246
	o2sat_imp1	1.01429313	9.925303e-01	1.0383737
A tibble: 10 × 4	hr_imp1	1.01168795	1.008049e+00	1.0153248
	temp_imp1	1.00784512	9.309287e-01	1.0914074
	sbp_imp1	1.00292389	9.978885e-01	1.0079298
	map_imp1	0.98340323	9.760605e-01	0.9908138
	resp_imp1	1.05165073	1.038685e+00	1.0646521
	factor(hospid)A	0.00410084	9.841101e-05	0.1554736
	factor(hospid)B	0.00267482	6.446135e-05	0.1010534

First observation carried backwards

```

[19]: model_imp2 <- glm(
  # the -1 term will suppress intercept, which is needed to avoid singularity
  ↪matrix issue
  any_sepsis ~ -1 + age + gender + o2sat_imp2 + hr_imp2 + temp_imp2
  + sbp_imp2 + map_imp2 + resp_imp2 + factor(hospid),

```



```
data = filter(df_imp, iculos == 1),
family = binomial(link = "logit")
)
```

```
[20]: summary(model_imp2)
```

Call:

```
glm(formula = any_sepsis ~ -1 + age + gender + o2sat_imp2 + hr_imp2 +
temp_imp2 + sbp_imp2 + map_imp2 + resp_imp2 + factor(hospid),
family = binomial(link = "logit"), data = filter(df_imp,
iculos == 1))
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.1188	-0.3487	-0.2969	-0.2530	2.9213

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
age	0.002011	0.001858	1.083	0.278993	
gender	0.179775	0.060351	2.979	0.002894	**
o2sat_imp2	-0.003505	0.008713	-0.402	0.687495	
hr_imp2	0.011096	0.001669	6.647	2.98e-11	***
temp_imp2	0.036431	0.038238	0.953	0.340720	
sbp_imp2	0.001067	0.002164	0.493	0.622103	
map_imp2	-0.012117	0.003177	-3.814	0.000137	***
resp_imp2	0.034078	0.005201	6.552	5.69e-11	***
factor(hospid)A	-4.629805	1.652233	-2.802	0.005076	**
factor(hospid)B	-5.080070	1.648675	-3.081	0.002061	**

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 33309.9 on 24028 degrees of freedom  
Residual deviance: 9339.9 on 24018 degrees of freedom  
(6897 observations deleted due to missingness)  
AIC: 9359.9

Number of Fisher Scoring iterations: 6

```
[21]: # Calculate odds ratio & 95% confidence interval
exp(coefficients(model_imp2)) %>%
  enframe(name = "variable", value = "odds ratio") %>%
  add_column(as_tibble(exp(confint(model_imp2))))
```

Waiting for profiling to be done...

	variable <chr>	odds ratio <dbl>	2.5 % <dbl>	97.5 % <dbl>
A tibble: 10 × 4	age	1.002013075	0.9983839947	1.0056812
	gender	1.196947746	1.0637620110	1.3477559
	o2sat_imp2	0.996501380	0.9804169530	1.0145076
	hr_imp2	1.011157302	1.0078466742	1.0144630
	temp_imp2	1.037102407	0.9623432758	1.1179221
	sbp_imp2	1.001067178	0.9968119421	1.0053017
	map_imp2	0.987955994	0.9818196556	0.9941161
	resp_imp2	1.034665510	1.0240921583	1.0452040
	factor(hospid)A	0.009756666	0.0003699473	0.2408788
	factor(hospid)B	0.006219472	0.0002375450	0.1525313

### 1.5.3 Higher respiration rate among those with sepsis?

```
[22]: df_imp %>%
  group_by(patid) %>%
  filter(iculos == 1) %>%
  ungroup() %>% group_by(any_sepsis) %>%
  skim(resp_imp1)
```

```
Data Summary
Name                               Values
Number of rows                    30925
Number of columns                  32
-----
Column type frequency:
  numeric                          1
-----
Group variables                    any_sepsis

Variable type: numeric

  skim_variable any_sepsis n_missing complete_rate mean    sd    p0    p25    p50
p75  p100 hist
1 resp_imp1      0         960          0.967  18.2  4.42    1   15.4   17.9
20.4  98
2 resp_imp1      1          68          0.959  19.6  5.54    1   15.9   18.8
22.4  44.5
```

```
[23]: df_imp %>%
  group_by(patid) %>%
  filter(iculos == 1) %>%
  ungroup() %>% group_by(any_sepsis) %>%
  skim(resp_imp2)
```

```

Data Summary
Name          Values
Number of rows Piped data
Number of columns 30925
Number of columns 32

-----
Column type frequency:
  numeric          1
-----
Group variables    any_sepsis

Variable type: numeric

skim_variable any_sepsis n_missing complete_rate mean    sd    p0    p25    p50
p75  p100 hist
1 resp_imp2          0      960          0.967  18.2  5.20    1    15    18
21      98
2 resp_imp2          1      68          0.959  19.5  6.08    1   15.5   19
22.5    50

```